CS F213 / MAC F212 · OBJECT-ORIENTED PROGRAMMING · SEMESTER I 2026–2027

Intelligent IT Service Management & Incident Resolution

I-ITSM — project report

A role-based incident lifecycle platform with explainable priority, configurable SLAs, technician assignment strategies, notifications, audit history, reporting and intelligent triage.

Spring Boot 4.1.2 | PostgreSQL 17 | React 19 | Vite 8 | TypeScript 7

Team
Complete names / roll numbers / responsibilities before submission

Submission
On or before 22 November 2026

Abstract

I-ITSM is a full-stack ITSM platform that manages the complete lifecycle of an IT incident: report, classify, prioritize, assign, acknowledge, investigate, resolve, verify and close. It supports three user groups — employees, technicians and administrators — with role-specific experiences and secure JWT-based access. The backend applies object-oriented design through a state machine for lifecycle transitions, strategy objects for priority calculation and technician assignment, factories, repositories, domain events and an append-only audit trail. PostgreSQL stores the operational data and Flyway owns the schema. The React frontend gives each role a purpose-built dashboard and workflow.

The system goes beyond a basic CRUD application with rule-based auto-classification, duplicate incident detection, smart technician recommendations, SLA warning/breach detection, automatic escalation and suggested resolutions from historical incidents. It was validated with 34 automated tests (including a real PostgreSQL Testcontainers integration test) and 60+ documented scenarios covering authorization, invalid input, lifecycle rules and boundary conditions.

Table of contents

1. Problem analysis
2. Requirements specification
3. System architecture
4. Object-oriented design and UML
5. Database design
6. Intelligent features
7. Security
8. Technology stack and justification
9. Testing and validation
10. Demonstration workflows
11. Evaluation rubric mapping
12. Limitations and future work
13. Conclusion
14. Appendices

1. Problem analysis

1.1 Problem

Large organisations depend on many IT services — payment platforms, identity providers, databases, networks, storage and collaboration tools. When a service fails, the employee who notices it has no structured way to report the problem, the support team must manually triage a flood of reports, and managers lack timely visibility into service health and SLA performance. Common failures include:

- inconsistent incident descriptions and missing affected-service information;
- manual and subjective prioritisation;
- delayed or uneven technician allocation;
- missed SLAs and no early warning;
- duplicate reports of the same underlying failure;
- lost incident history and weak accountability;
- reports that cannot answer operational questions such as “which service is failing most?”.

1.2 Stakeholders

STAKEHOLDER	GOALS	PAIN POINTS ADDRESSED
Employee / customer	Report an issue quickly, track progress, provide extra information, verify the fix and give feedback	Simple report form, live status, comments, feedback
Technician	See assigned work, understand urgency, acknowledge, investigate, request information, resolve and manage workload	Personal queue, SLA clocks, internal notes, workload capacity
Administrator / IT manager	Manage people and services, configure SLAs and assignment, monitor incidents, review reports and audit	User/catalogue/SLA/settings admin, dashboards, audit trail
Organisation	Reduce downtime, meet contractual targets, make evidence-based staffing decisions	Priority automation, SLA monitoring, workload and service analytics

1.3 Existing challenges

1. **Triage does not scale.** Manually reading every incident delays assignment and makes priority inconsistent.
2. **Assignment is opaque.** Choosing a technician without considering skills, load and experience leads to burnout and slow resolution.
3. **SLAs are reactive.** Deadlines are often discovered only after they are missed.
4. **Duplicates waste effort.** Similar reports arrive independently and are investigated separately.
5. **History is fragmented.** Chat threads and spreadsheets do not provide an auditable lifecycle.
6. **No role boundaries.** Employees, technicians and managers need different access and different views; a shared spreadsheet cannot provide that safely.

1.4 Proposed solution

I-ITSM provides one integrated platform with:

- a guided employee report form with automatic category suggestions and duplicate checks;

- explainable, rule-based priority scoring from impact, urgency, service criticality and affected users;
- configurable SLA policies per priority and per service, with response/resolution clocks, warnings, breach flags and escalation;
- three interchangeable technician assignment strategies plus ranked recommendations;
- in-app and simulated-email notifications driven by domain events;
- a complete, append-only audit trail;
- searchable, paginated incident queues and role-scoped dashboards/reports;
- administration screens for users, skills, services, categories, SLA policies and routing settings.

1.5 Assumptions

- Email/SMS delivery is simulated in the course environment; the notification abstraction allows a real provider to be added later.
- SLA targets in the demo are in minutes so that warnings and breaches can be observed during a short demonstration.
- The deployment is a single-tenant service desk. The model could be extended with a tenant column if required.
- Administrators are trusted service managers; privileged operations are still audited.
- The demo seeder is disabled in a production profile.

---

2. Requirements specification

2.1 User roles

ROLE	CAPABILITIES
Employee	Register/login, report incidents, view own incidents, add information, verify resolution, provide feedback, receive notifications
Technician	Login, view assigned incidents, acknowledge, start investigation, add internal/public notes, resolve, view workload, receive notifications
Administrator	Manage users, technicians and skills; manage services/categories; configure SLAs; choose assignment strategy; monitor incidents and SLA breaches; view reports and audit logs; assign/reassign; override priority; close/cancel

2.2 Functional requirements

FR-1 User management and authentication

- FR-1.1 Employees can self-register; administrators can create all roles.
- FR-1.2 Passwords are hashed; login returns a signed JWT.
- FR-1.3 Every endpoint is protected by authentication except login/registration/health.
- FR-1.4 Role checks and per-incident ownership checks prevent unauthorized access.
- FR-1.5 Administrators can deactivate users, update technician capacity/skills and reset passwords.

FR-2 Incident management

- FR-2.1 Employees can create incidents with title, description, service, impact, urgency and affected users.
- FR-2.2 The system auto-classifies the category when none is supplied.
- FR-2.3 The system calculates priority and attaches the applicable SLA.
- FR-2.4 Employees can edit an incident before assignment; edits re-run classification.
- FR-2.5 Administrators can assign/reassign, override priority, cancel early incidents and close verified incidents.
- FR-2.6 Technicians can acknowledge, start and resolve assigned incidents with notes.
- FR-2.7 Reporters can verify a resolution, reopen a resolved incident and leave feedback.
- FR-2.8 The system rejects invalid lifecycle transitions.

FR-3 Priority management

- FR-3.1 Priority uses impact, urgency, service business criticality and affected users.
- FR-3.2 Every priority includes a score and a human-readable rationale.
- FR-3.3 P1/P2/P3/P4 map to configurable SLA policies.

FR-4 Assignment

- FR-4.1 At least two assignment strategies are available; I-ITSM implements three.
- FR-4.2 The strategy is configurable at runtime by an administrator.
- FR-4.3 The system scores skill match, load, experience and service history.
- FR-4.4 Technicians at capacity or unavailable are not auto-assigned.
- FR-4.5 Administrators can see recommendations and assign manually.

FR-5 SLA management

- FR-5.1 Policies define response/resolution targets and a warning percentage.
- FR-5.2 Policies can be global or service-specific; service-specific wins.
- FR-5.3 Response and resolution deadlines are calculated from creation time.
- FR-5.4 A scheduled monitor records warnings, breaches and escalation.
- FR-5.5 SLA status is exposed in list/detail APIs and dashboards.

FR-6 Notifications

- FR-6.1 In-app notifications are persisted and readable/markable.
- FR-6.2 Simulated email notifications are persisted and logged.
- FR-6.3 Creation, assignment, status change, comments, SLA warning/breach, escalation and duplicate detection produce notifications.
- FR-6.4 New channels can be added by extending the event listener without changing the domain.

FR-7 Audit

- FR-7.1 Every significant action records actor, action, previous/new state, description and time.
- FR-7.2 The audit trail is append-only and visible per incident and globally to administrators.

FR-8 Search, filtering and reporting

- FR-8.1 Search by title/description/reference.
- FR-8.2 Filter by status, priority, category, service, technician, reporter, SLA state, date range, duplicates and escalations.
- FR-8.3 Paginated result sets with deterministic ordering.
- FR-8.4 Dashboards show totals, priority/status mix, category/service analysis, seven-day trend, recent incidents, averages and technician workload.

FR-9 Intelligent features

- FR-9.1 Keyword-based auto-classification with confidence and rationale.
- FR-9.2 Duplicate detection with similarity scores and linking.
- FR-9.3 Ranked technician recommendations.
- FR-9.4 Suggested resolutions from similar resolved incidents.
- FR-9.5 SLA risk detection and automatic escalation.

FR-10 Administration

- FR-10.1 CRUD-like management (create/update/deactivate) for users, services and categories.
- FR-10.2 SLA policy management with validation.
- FR-10.3 Assignment strategy and auto-assign configuration.
- FR-10.4 Global audit view.

2.3 Non-functional requirements

REQUIREMENT	IMPLEMENTATION
Security	JWT, BCrypt, RBAC, ownership checks, validation, CORS, ProblemDetail
Maintainability	Feature packages, interfaces/strategies, DTO mapping, Flyway schema

REQUIREMENT	IMPLEMENTATION
Testability	Pure unit-testable rules + Testcontainers integration test
Performance	Indexed filters, pagination, entity graphs, batched lazy loading
Reliability	Transactions, optimistic locking, idempotent SLA monitor
Usability	Role-specific dashboards, guided forms, visible SLA state, responsive UI
Portability	Dockerised PostgreSQL; environment-driven config
Observability	Actuator health, audit trail, structured logs with simulated email
Extensibility	New categories/services/SLA policies from admin; new strategies/notifications by adding beans/listeners
Data integrity	Constraints, unique policies, reference uniqueness, validation

2.4 Business rules

1. An incident cannot jump directly from NEW to CLOSED; the lifecycle must be followed.
2. Only the reporter or an administrator can edit an incident, and only before assignment.
3. Only the assigned technician can acknowledge/start/resolve.
4. Only the reporter can verify and submit feedback.
5. Only an admin can close a verified incident, cancel early incidents, reassign or override priority.
6. Feedback is allowed after resolution, exactly once.
7. A resolution target must be longer than its response target.
8. A technician cannot be assigned beyond capacity or while unavailable/inactive.
9. Deactivating a technician with active incidents is blocked; the last active admin cannot be deactivated; an admin cannot deactivate themselves.
10. Service-specific SLA policies override global policies of the same priority.
11. Internal notes are visible only to technicians and administrators.
12. P1/P2 incidents not acknowledged by the response deadline escalate over time.

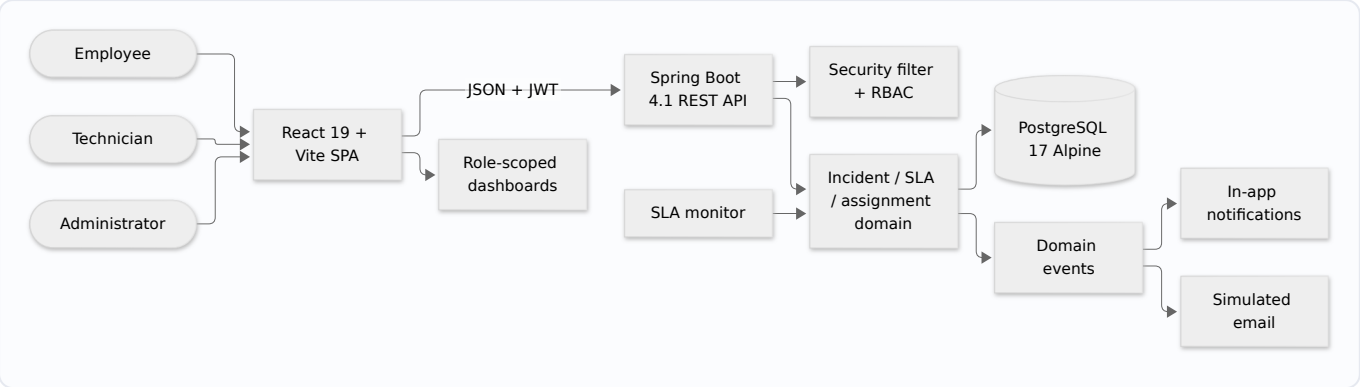
3. System architecture

3.1 Architectural style

I-ITSM is a **modular monolith** with a REST API and a separate SPA. A monolith was chosen because:

- the domain is cohesive and the expected load is moderate;
- a single transaction can keep incident state, audit and notifications consistent;
- it is easier to deploy and demonstrate within a course timeline;
- package boundaries and interfaces keep the door open for extracting services later.

3.2 High-level architecture



The SPA uses a typed API client and a Vite development proxy (`/api` , `/actuator`) to avoid CORS complexity during development. The backend is stateless; the JWT carries identity and roles.

3.3 Backend components

COMPONENT	RESPONSIBILITY
Security	JWT filter, authentication provider, current-user service, security error handlers
User	Inheritance hierarchy, authentication, admin user management, skills
Catalogue	Services with criticality, categories, activation
SLA	Policy resolution, deadline arithmetic, status resolver, scheduled monitor, escalation
Incident	Aggregate, lifecycle orchestration, search, DTOs, comments, access control
Lifecycle	State pattern classes and transition validation
Priority	Weighted priority strategy and explainable result
Assignment	Candidate scoring, three strategies, recommendations
Intelligence	Classification, duplicate similarity, suggested resolutions
Audit	Append-only event log
Notification	Domain events, observers, in-app and simulated email
Dashboard	Role-scoped aggregation and analytics

COMPONENT	RESPONSIBILITY
Settings	Runtime assignment strategy/auto-assign
Seed	Realistic demonstration data

3.4 Frontend structure

- A role-aware navigation shell and protected routes.
- Dashboard with KPI cards, Recharts trend/donut/bar charts and SLA watch.
- Incident queue with search + six visible filters and pagination.
- Report form with live smart-triage analysis and duplicate preview.
- Incident detail with SLA clocks, metadata, duplicate warning, suggestions, comments/internal notes, audit timeline and role-specific actions.
- Technician queue with capacity utilisation and deadline ordering.
- Admin screens for users/skills, services/categories, SLA policies, assignment settings and audit.
- Notification inbox with channel filters and unread handling.
- Responsive layout, mobile navigation drawer and accessible controls.

3.5 Request flow and consistency

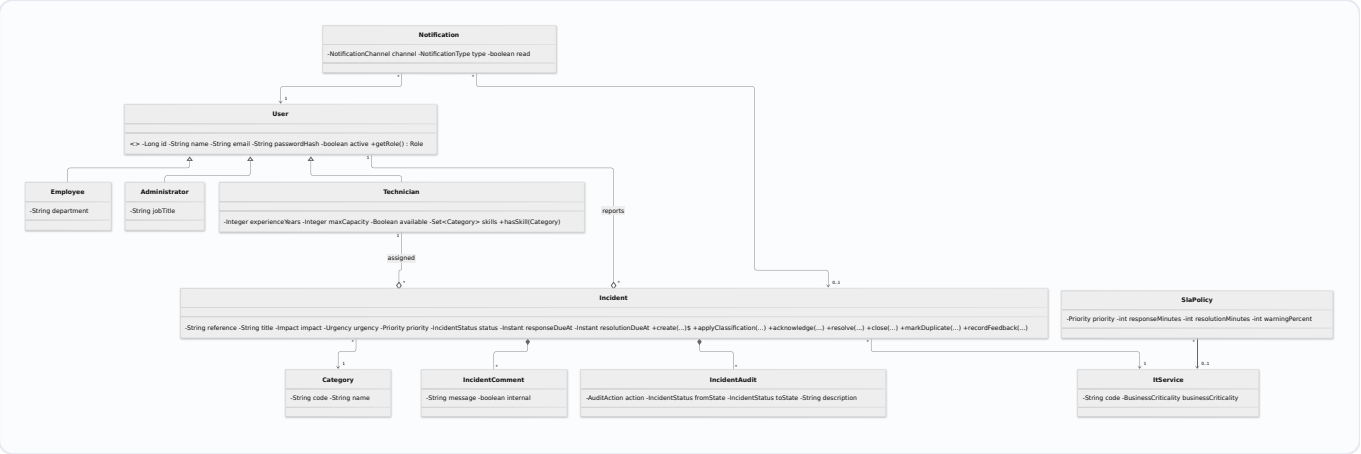
Incident creation, transitions, comments and admin changes run in a single transaction. Domain events are synchronous, so notification and audit writes commit with the incident change. Incidents use optimistic locking (`@Version`) to prevent lost updates when two actors change the same incident. SLA monitor operations are idempotent through persisted warning/breach flags.

---

4. Object-oriented design and UML

4.1 Class diagram (simplified)

Domain classes

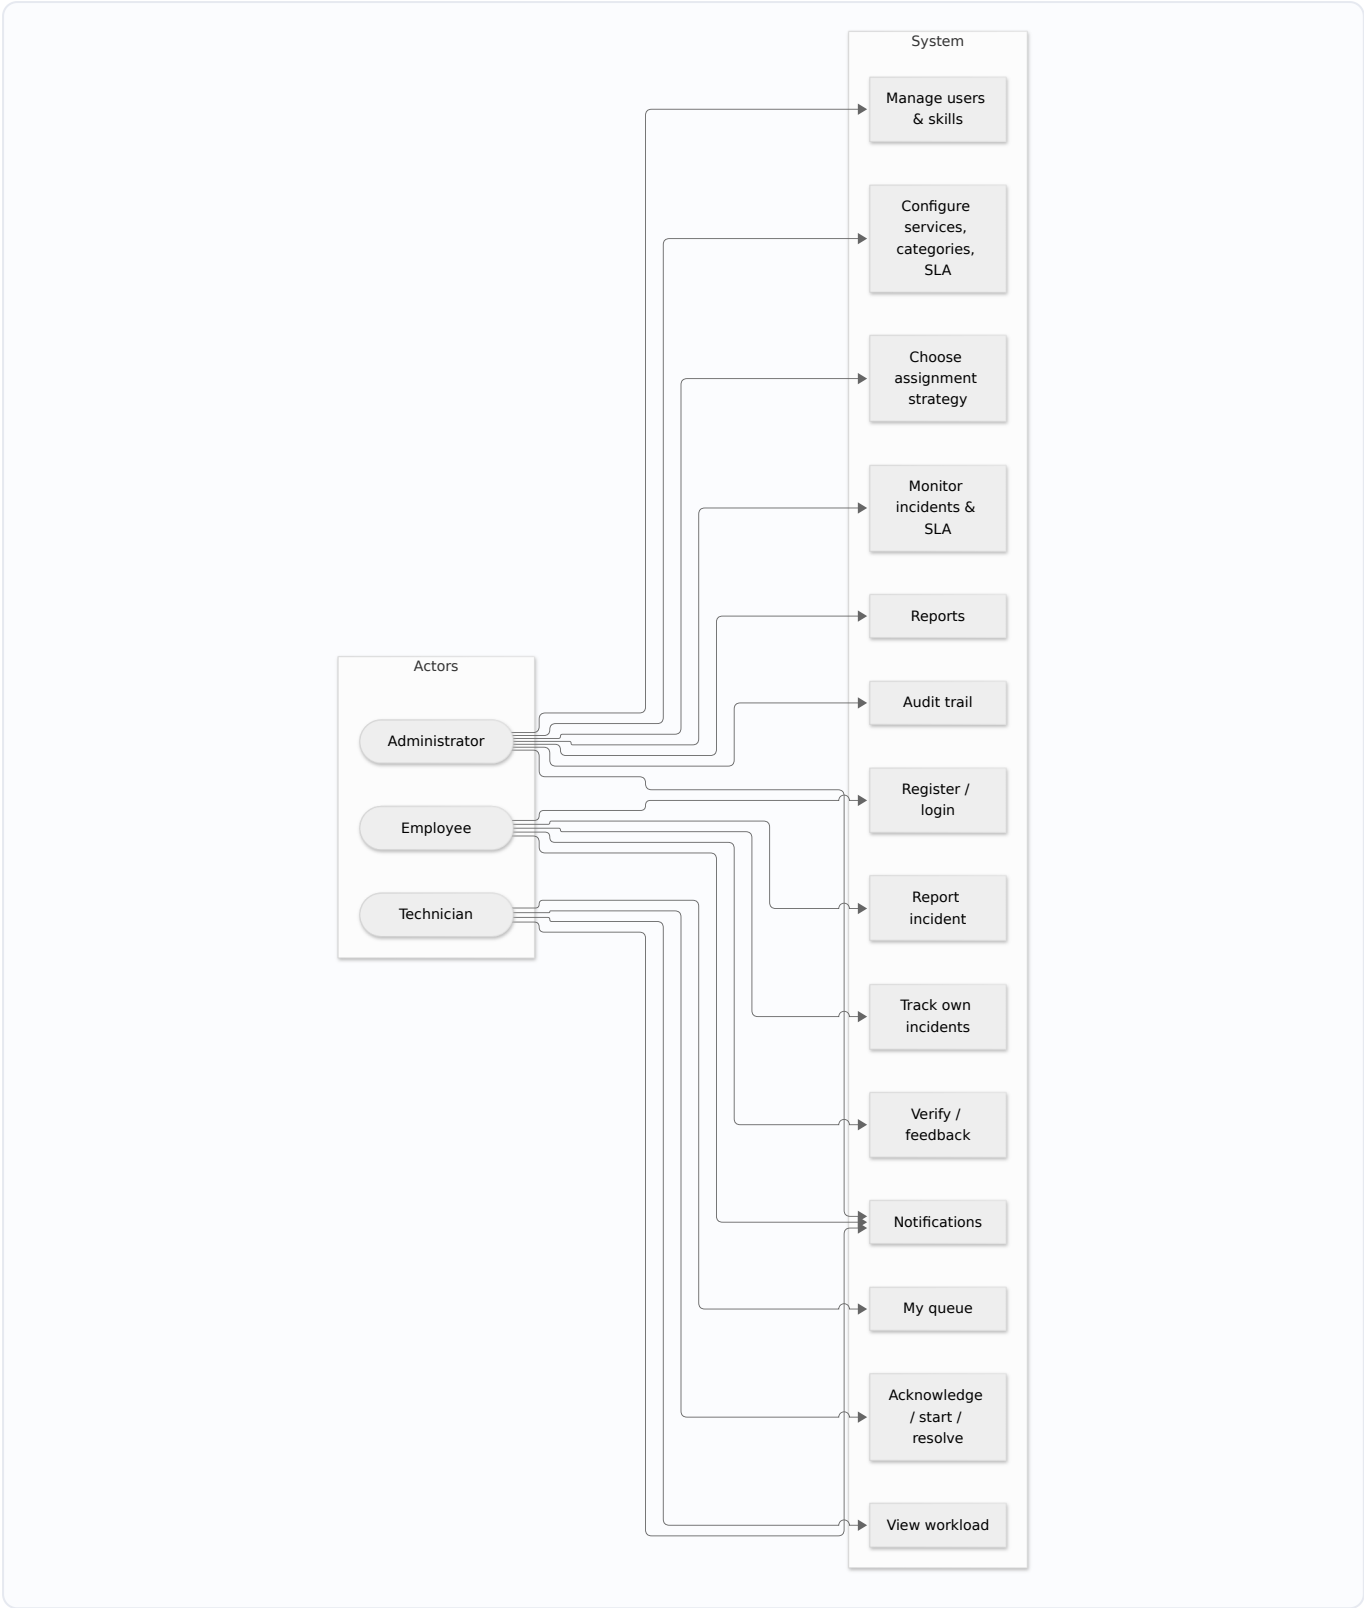


Strategy and state contracts

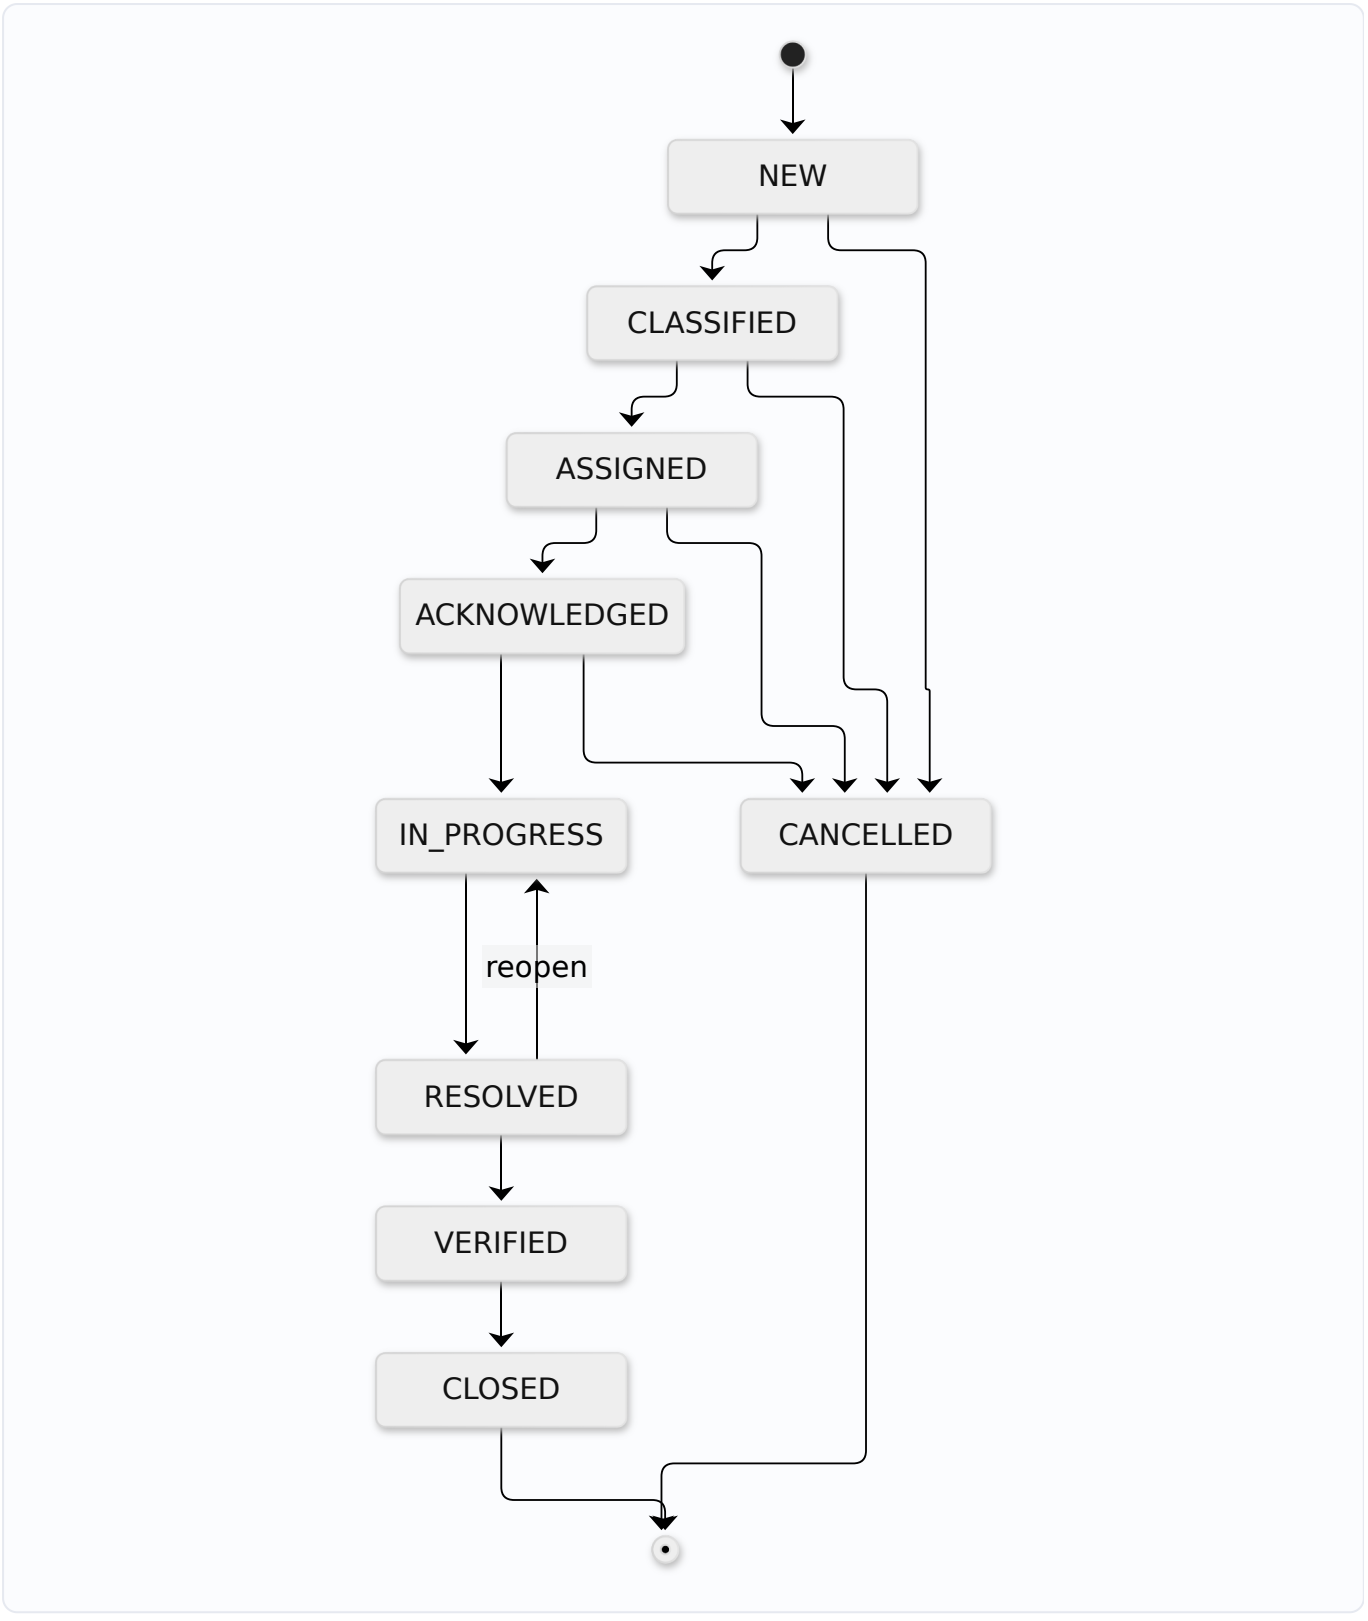


One concrete `IncidentState` class exists for each status: `NEW`, `CLASSIFIED`, `ASSIGNED`, `ACKNOWLEDGED`, `IN_PROGRESS`, `RESOLVED`, `VERIFIED`, `CLOSED` and `CANCELLED`. The state diagram in 4.3 shows the transition graph. The full diagram with every state class and attribute is in ARCHITECTURE.md.

4.2 Use-case diagram



4.3 State diagram



4.4 How OOP concepts are demonstrated

CONCEPT	EVIDENCE
Encapsulation	Private entity fields; entity behaviour methods (`acknowledge` , `resolve` , `markDuplicate`); services enforce access; controllers only orchestrate

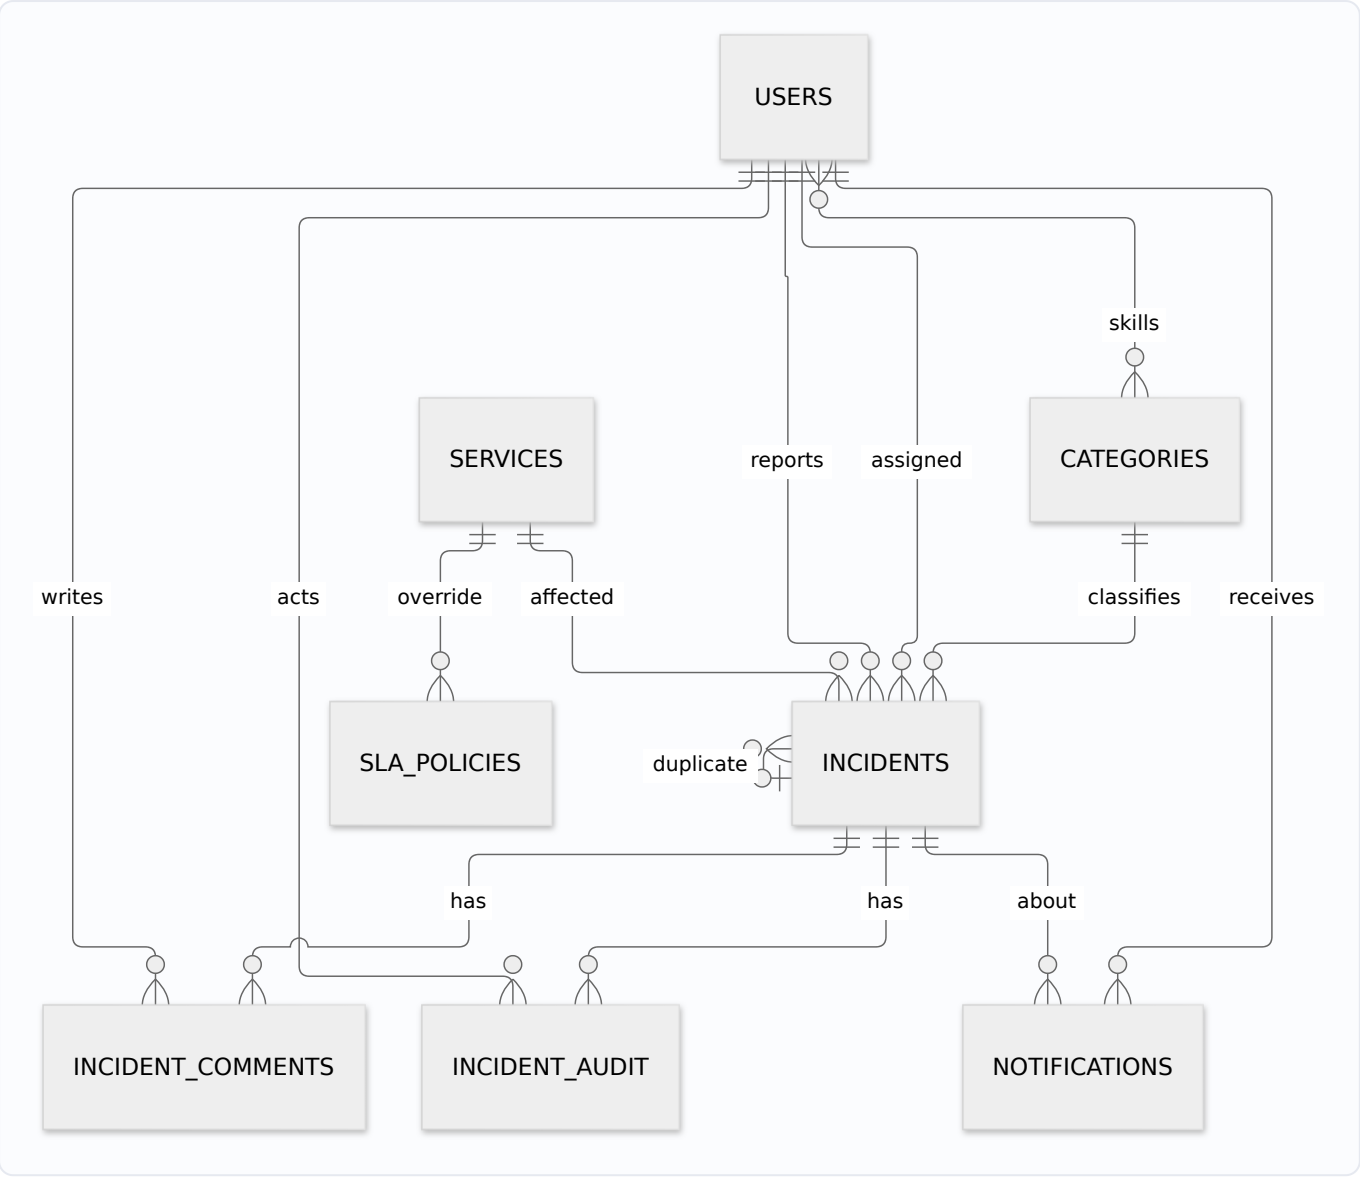
CONCEPT	EVIDENCE
Abstraction	`PriorityStrategy`, `AssignmentStrategy`, `IncidentState`, repository interfaces
Inheritance	`User` hierarchy; `AbstractIncidentState`; concrete strategy classes implement shared contracts
Polymorphism	`getRole()` overridden per subtype; runtime strategy selection; state classes selected per status
Interfaces	Strategy/state/repository contracts, Spring Data repositories
Composition	`Incident` owns comments and audit records; services compose collaborators
Association/aggregation	Incident → reporter/technician/category/service; technician → skills
Exception handling	Domain exceptions mapped to HTTP statuses with ProblemDetail; state transition and capacity errors
Collections	Skill sets, candidate lists, audit timelines, paginated content, filtered queries

4.5 Design patterns

PATTERN	IMPLEMENTATION	BENEFIT
State	`IncidentState` + concrete states	Illegal transitions rejected in one place
Strategy	Priority + three assignment strategies	Change algorithms without editing the service
Factory	User, notification, strategy, state factories	Centralized creation
Observer	Spring domain events + notification listener	Notifications/audit decoupled from workflow
Repository	Spring Data JPA	Persistence abstraction and testability
Builder	Demo incident seeder	Complex optional lifecycle data assembled readably
Adapter	Simulated email channel	Real provider can replace it later
Template method	Abstract state class	Reuses status/transition storage
DTO/mapper	Request/response records + mapper	Stable API contracts, no entity leakage

5. Database design

I-ITSM uses PostgreSQL 17 with Flyway migrations. The complete ER diagram and column-level schema are in DATABASE.md. The core relationships are:



Key design choices:

- **Single-table user inheritance** with a `user_type` discriminator. Authentication queries stay simple and role-specific columns are nullable for other subtypes.
- **Normalized catalogue**: services and categories are independent aggregates referenced by incidents and technician skills.
- **SLA policy scoping**: global policies use a partial unique index on priority; service-specific policies use a unique `(priority, service_id)` index.
- **Incident history** is normalized into comments and audit tables; both cascade on incident deletion, while audit history remains append-only through the API.
- **Duplicate relationship** is a self-reference with an optional `ON DELETE SET NULL`.
- **Optimistic locking** through `version`.
- **Indexes** support status, priority, reporter, technician, created-at and resolution-due queries.

Sample data is generated by the seeder: 7 users, 4 technicians, 10 categories, 8 services, 5 SLA policies (global + service-specific) and 13 incidents across every lifecycle stage, with comments, audit records, notifications, breaches, escalation and feedback.

---

6. Intelligent features

6.1 Explainable priority

```
raw   = impact(1–4) × urgency(1–4) × business criticality(1–4)
bonus = 0 / 1 / 3 / 6 / 10 for affected users <5 / ≥5 / ≥25 / ≥100 / ≥500
score = raw + bonus
P1 ≥ 36 · P2 ≥ 18 · P3 ≥ 6 · P4 otherwise
```

A total outage (`EXTREME` impact) of a `CRITICAL` service is forced to P1. The rationale string is stored on the incident and shown in the UI, so users and evaluators can see exactly why a priority was assigned. Administrators can override the priority with a reason; the override is audited and recalculates SLA deadlines.

6.2 Auto-classification

An active-category keyword map scores the title and description (title matches count double). The best category is returned with a confidence and rationale. For example, “VPN disconnects / Wi-Fi / DNS” classifies as Network, “password / SSO / MFA” as Authentication, and unmatched reports fall back to Other for manual triage. New categories added by an administrator are immediately available to the classifier.

6.3 Duplicate detection and correlation

Text tokenization removes stop words and compares the title and description using an overlap coefficient. A same-service bonus increases the score. Incidents above a candidate threshold are returned as related incidents; strong matches are linked as possible duplicates (`duplicate_of_id` , `duplicate_score`) and administrators are notified. The UI displays the duplicate banner and links to the related reports. This directly addresses the brief's duplicate-report problem.

6.4 Smart technician recommendation

Each candidate receives a transparent score:

- 40 points for a category skill match;
- up to 30 points for free capacity;
- up to 15 points for experience;
- up to 15 points for successful history on the affected service;
- penalties for unavailable/at-capacity technicians.

The assignment strategies consume the same candidate set, while the recommendation endpoint shows the ranked list with load, capacity, skills and service history for manual override.

6.5 SLA risk, breaches and escalation

The monitor scans unresolved incidents for warning thresholds, response breaches, resolution breaches and unacknowledged P1/P2 incidents. Flags make it idempotent. An at-risk/breached/on-track/met SLA state is computed from deadlines and flags and is exposed in every incident response and dashboard. Escalation levels increase over time for unacknowledged critical incidents, notifying administrators and adding audit entries.

6.6 Suggested resolution

Resolved incidents from the last 180 days are compared with the current incident using the same similarity engine. The top matches return the resolution notes, timestamps and match scores, giving technicians a starting point and improving organisational knowledge reuse.

7. Security

Implemented controls:

- BCrypt password hashing; password hashes never serialized.
- HS384 JWT with expiration; invalid/tampered tokens rejected.
- All APIs except login/registration/health require authentication.
- Method-level role checks plus service-level ownership checks:
  - employees only see incidents they reported;
  - technicians only see incidents assigned to them;
  - admins have organisation-wide access;
  - internal notes are filtered from employee responses.
- Bean Validation on request DTOs, maximum lengths and enum values.
- RFC 9457 ProblemDetail responses; no stack traces or database details leaked.
- CORS allow-list; configurable origins.
- Optimistic locking; self-deactivation and last-admin guards; active-work deactivation guard.
- Demo accounts and seeded passwords are clearly development-only.

Limitations (documented for production): no refresh-token rotation/revocation, no rate limiting/ account lockout, no HTTPS termination in the development setup, simulated email, no MFA, and audit rows are immutable by convention/API rather than by a database trigger. The report's future-work section lists these.

8. Technology stack and justification

LAYER	TECHNOLOGY	VERSION	ROLE
Language	Java	21 target (JDK 25 used to build/test)	Strong typing, mature OOP ecosystem, Spring support
Backend	Spring Boot	4.1.2	Web MVC, dependency injection, security, data access, scheduling
Security	Spring Security + JJWT	7.x / 0.12.6	JWT authentication, BCrypt, role-based authorization
Persistence	Spring Data JPA + Hibernate	Boot-managed	Repository abstraction, derived queries, specifications, optimistic locking
Migrations	Flyway	Boot-managed	Versioned schema evolution
Database	PostgreSQL	17 Alpine (Docker)	Relational integrity, partial indexes, timestamptz, mature tooling
Frontend	React	19.3	Component model, large ecosystem, role-specific interfaces
Build/dev	Vite	8.3	Fast dev server, proxy, TypeScript bundling
Language (UI)	TypeScript	7.0	Typed API contracts, safer refactors
Routing	React Router	7.x	Protected role-aware routes
Charts	Recharts	3.x	Responsive dashboards
Testing	JUnit 6, AssertJ, Mockito, Testcontainers, Playwright	Boot-managed / current	Unit, integration, browser validation
Infra	Docker Compose	current	Reproducible PostgreSQL with healthcheck/volume

8.1 Why Java?

Java offers a mature OOP runtime, static typing, excellent collections and a rich Spring ecosystem. The project explicitly needs interfaces, inheritance, exceptions and a maintainable layered design; Java expresses these clearly, and Spring Boot removes infrastructure boilerplate while keeping the domain testable. Type safety also makes large domain refactors safer.

8.2 Why React?

React's component model maps naturally to role-specific pages and reusable UI (badges, SLA clocks, tables, modals). Its ecosystem provides routing, charts and icons, and Vite gives fast feedback. The frontend is a separate deployable SPA, which keeps the API language-independent.

8.3 Why Spring Boot?

Spring Boot provides production-grade REST, security, JPA, scheduling, validation and actuator support with minimal configuration. Dependency injection supports the strategy/factory/observer designs, method security expresses role rules cleanly, and the test slices/Testcontainers support make both unit and integration testing practical.

8.4 Why PostgreSQL?

Incidents, users, audits and SLA policies are highly relational and benefit from foreign keys, transactions and constraints. PostgreSQL adds partial unique indexes for SLA scoping, `timestamptz` for deadlines, powerful filtering and a trustworthy concurrency model. It is also free, widely available and easy to run in Docker Alpine. A document database would make reporting and relational integrity harder for this domain.

8.5 Why a modular monolith + SPA?

The domain is cohesive and transaction consistency matters (incident + audit + notification). A single backend is easier to deploy and demonstrate than microservices, while feature packages and interfaces provide clear boundaries. The SPA provides a richer, role-specific UX than server-rendered pages and can be hosted independently.

8.6 Alternatives considered

- **Node/Express or Django:** viable, but weaker fit for the requested OOP emphasis and the team's Java/Spring experience.
- **MongoDB:** flexible but poorer for the reporting joins, referential integrity and SLA policy constraints.
- **Microservices:** overkill for the project size and would complicate transactions and local deployment.
- **Server-rendered UI:** simpler, but less suitable for the interactive dashboards, timeline and role-specific workflows.
- **Client-side database:** rejected because persistence and multi-user consistency are core requirements.

8.7 Advantages and limitations of the stack

Advantages: strong typing, mature security, excellent relational support, fast frontend feedback, easy local infrastructure, and a large ecosystem.

Limitations: Spring's startup time and memory footprint; JPA can hide N+1 queries if misused; the SPA adds a separate build/deploy step; JWT revocation and refresh rotation are non-trivial; Docker is required for the integration test. These are manageable at course-project scale.

8.8 How the stack supports OOP design

Java interfaces/abstract classes directly model strategies and states; Spring's DI makes them swappable; JPA repositories express persistence contracts; exception hierarchies map domain failures cleanly; React components and TypeScript interfaces mirror the domain DTOs; and the separation between controllers, services, aggregates and repositories keeps responsibilities clear.

8.9 Evolution if users increase significantly

Add database-level reporting aggregates/materialized views, move SLA scanning and email into a queue/worker, introduce refresh tokens and rate limiting, partition or archive old incidents, cache role dashboards, add OpenTelemetry metrics/tracing, and extract notification/reporting services if scale demands it. The existing strategy/state/event boundaries make these changes incremental.

---

9. Testing and validation

Automated result: 34/34 tests pass (33 unit + 1 integration). The integration test uses a real `postgres:17-alpine` Testcontainers instance and exercises the complete HTTP workflow.

AREA	AUTOMATED TESTS	MANUAL VERIFICATION
Security	JWT validation/tamper; auth integration	Role boundaries, invalid login, 401/403, self-deactivation, internal notes
Lifecycle	State transitions; integration happy path	Edit/reopen/cancel/illegal transitions
Priority/SLA	Formula, bands, deadlines, statuses	Service-specific SLA, dashboard SLA states, monitor scan
Assignment	Three strategies, capacity filtering	Live strategy switch, recommendations
Intelligence	Classification, similarity	Duplicate preview/links, suggested resolutions
Reporting	—	Role-scoped dashboards, reports, filters, pagination
Audit/notifications	—	Full workflow timeline, channel rows, read/unread
Persistence	Flyway/Testcontainers	Fresh seed, reset and restart

TEST_CASES.md documents 60+ test cases with preconditions, steps, expected results and evidence. Additional negative checks were run manually for blank titles, unknown services, feedback before resolution, duplicate feedback, editing after assignment, closing before verification, deactivating a technician with active work, unknown endpoints and SLA monitor idempotence.

The UI was validated at 1440 px desktop and 390 px mobile widths with headless Chromium; no console or page errors were observed during login, dashboards, queue, detail, forms, admin pages and the full technician action flow.

10. Demonstration workflows

The demo plan is in DEMO_SCRIPT.md. The three required workflows are:

1. **Report → triage → assignment → notification** Employee reports a payment outage; the system auto-classifies it, calculates P1, attaches the payment-specific SLA, detects a possible duplicate, assigns the best technician and notifies both parties.
2. **Investigate → resolve → verify → close** Technician acknowledges, starts work, adds an internal note, resolves with notes; the reporter verifies; the admin closes; feedback is captured; the audit timeline and notifications show every step.
3. **Correlation → SLA risk → escalation → reporting** A similar payment report is linked as a duplicate; a breached/at-risk incident is visible on the dashboard; the SLA monitor escalates an unacknowledged critical incident; the reports page shows workload, service/category analysis and SLA health; the audit log shows the full history.

The 15–20 minute video plan assigns each team member a module and a workflow segment.

11. Evaluation rubric mapping

RUBRIC CRITERION	MARKS	EVIDENCE
System design & architecture	10	Modular architecture, state/strategy/observer/factory/repository patterns, documented diagrams
Functionality implementation (MVP)	60	All mandatory modules: users, incidents, priority, assignment, SLA, notifications, audit, search/reporting, intelligence, persistence
User experience	10	Role-specific responsive UI, dashboards, smart triage form, SLA clocks, admin screens
Technical depth	15	Spring Security/JWT, JPA specifications, Testcontainers, event-driven notifications, optimistic locking, Flyway, reusable typed React architecture
Testing & validation	5	34 automated tests, 60+ documented cases, negative/boundary checks
Innovation	10	Auto-classification, duplicate detection, technician recommendations, suggested resolutions, SLA risk detection, auto-escalation
Documentation & presentation	10	README, architecture/UML, ER design, API reference, test catalogue, report and demo script

12. Limitations and future work

Current limitations

- Email/SMS are simulated (stored + logged), not delivered by an external provider.
- No mobile push notifications or real-time WebSocket updates; the bell polls every 30 seconds.
- JWT access tokens only; no refresh rotation/revocation list.
- No MFA or password-reset email flow.
- The dashboard aggregates in memory, suitable for the demo dataset but not millions of incidents.
- The SLA monitor runs inside the API process.
- No committed browser E2E suite or CI pipeline.
- Single-tenant; no organisation boundaries.

Future work

1. Notification adapters for SMTP/SMS/push/chat with an outbox and retries.
2. Real-time updates via SSE/WebSocket and a notification preference centre.
3. Refresh tokens, revocation, MFA, rate limiting and account lockout.
4. Database-level analytics, materialized views or a reporting warehouse for large volumes.
5. ML-assisted classification/similarity with human feedback loops, while keeping the rule-based fallback explainable.
6. Service dependency graph and incident correlation across related services.
7. Automated escalation policies configurable per service/priority, with on-call rosters.
8. CI/CD, Playwright E2E tests, load testing and observability dashboards.
9. Multi-tenancy and per-tenant SLAs.
10. Attachments, knowledge-base articles and a customer-facing portal.

13. Conclusion

I-ITSM demonstrates a complete, usable IT service management workflow rather than an isolated CRUD exercise. The design applies object-oriented principles to the parts of the problem where they matter: the incident lifecycle is a state machine, priority and assignment are interchangeable strategies, notifications and audit are observers of domain events, and entities protect their own invariants. The system meets every mandatory MVP requirement and adds meaningful intelligent features, a polished role-specific interface and a substantial automated test suite. Its documentation covers architecture, UML, database design, API contracts, tests and a video demonstration plan, providing a strong basis for evaluation and future extension.

14. Appendices

A. Demo accounts

ROLE	EMAIL	PASSWORD
Admin	`admin@itsm.local`	`Admin@123`
Technician	`priya@itsm.local`	`Tech@123`
Technician	`rohan@itsm.local`	`Tech@123`
Technician	`dev@itsm.local`	`Tech@123`
Employee	`employee@itsm.local`	`Employee@123`
Employee	`vikram@itsm.local`	`Employee@123`

B. Run commands

```
docker compose up -d --wait db
cd backend && ./mvnw spring-boot:run
# another terminal
cd frontend && npm install && npm run dev
```

If port 5432 is occupied:

```
POSTGRES_PORT=5433 docker compose up -d --wait db
cd backend && DB_URL=jdbc:postgresql://localhost:5433/app ./mvnw spring-boot:run
```

C. Reference documents

- README
- Architecture
- Database design
- API reference
- Test cases
- Demo script